Q: Which Azure AI workload is designed to extract insights from large volumes of unstructured data (like PDFs and images)?

A: Knowledge Mining. WHY: Combines Azure Cognitive Search with AI enrichment to make unstructured data searchable. Memory Cue: 'Mining = digging insights out of a mountain of data.' See also: Card 5 – Cognitive Search vs Knowledge Mining.

Q: A bank wants to monitor credit card transactions for unusual behavior. Which AI workload should they use?

A: Anomaly Detection. WHY: Identifies outliers and deviations in time-series or event-based data. Memory Cue: 'Anomaly = abnormal blip in the pattern.' See also: Card 7 – Use cases for Anomaly Detection.

Q: A business needs to extract key phrases and sentiment from customer reviews. Which AI workload applies?

A: Natural Language Processing (NLP). WHY: NLP allows systems to analyze and understand human language for tasks like sentiment analysis and key phrase extraction. Memory Cue: 'NLP = makes sense of words.' See also: Card 12 – Language Service features.

Q: Which AI workload is used for visual tasks like classifying objects in photos or detecting faces?

A: Computer Vision. WHY: Enables AI to process and analyze images for tagging, object detection, OCR, and more. Memory Cue: 'Vision = seeing and understanding images.' See also: Card 14 – Computer Vision vs Custom Vision.

Q: Which Responsible AI principle ensures systems treat all users fairly, avoiding discrimination in decisions like loan approvals?
A: Fairness. WHY: Ensures AI models do not propagate or amplify bias based on gender, ethnicity, or other factors. Memory Cue: 'Fair play for all.' See also: Card 20 – Fairness vs Inclusiveness.
Q: A healthcare provider wants to ensure AI systems protect patient confidentiality. Which principle applies?
A: Privacy and Security. WHY: Protects personal data and ensures ethical data use. Memory Cue: 'Privacy = protect the person.' See also: Card 22 – Privacy in healthcare AI.

Q: A company adds in-depth documentation so users understand how their AI system works and its limitations. Which principle is this?
A: Transparency. WHY: Makes AI systems understandable and their purpose and constraints clear. Memory Cue: 'Transparency = nothing hidden.' See also: Card 25 – Transparency vs Accountability.
Q: An AI startup sets up strict oversight committees to monitor AI deployments for ethical and legal compliance. Which principle?
A: Accountability. WHY: Ensures governance structures are in place to make AI traceable and answerable. Memory Cue: 'Accountable AI = someone’s responsible.' See also: Card 28 – Governance frameworks.

Q: An AI system for public services is tested to ensure it benefits all demographics equally. Which principle applies?

A: Inclusiveness. WHY: Ensures AI benefits reach all users, regardless of background. Memory Cue: 'Inclusiveness = all welcome.' See also: Card 20 – Fairness vs Inclusiveness.

Q: A company uses rigorous testing to ensure its autonomous vehicles don’t malfunction in unpredictable conditions. Which principle?

A: Reliability and Safety. WHY: AI must function consistently under expected and unexpected scenarios. Memory Cue: 'Reliable = safe to trust.' See also: Card 30 – Safety testing in AI systems.

Q: How is Fairness different from Inclusiveness in Responsible AI?

A: Fairness focuses on removing bias from AI outputs (e.g., loan approvals), while Inclusiveness ensures AI solutions benefit all societal groups. Memory Cue: 'Fair = no bias, Inclusive = everyone gains.' See also: Card 5 – Responsible AI overview.

Q: How is Transparency different from Accountability?

A: Transparency ensures users understand how an AI works. Accountability ensures there is governance and oversight for AI’s operation. Memory Cue: 'Transparency = clarity, Accountability = responsibility.' See also: Card 28 – Governance frameworks.