Front (Question)	Back (Answer + WHY + Memory Cue + Cross-Ref)
Which Azure AI workload is designed to extract insights from large volumes of unstructured data (like PDFs and images)?	Knowledge Mining. WHY: Combines Azure Cognitive Search with AI enrichment to make unstructured data searchable. Memory Cue: 'Mining = digging insights out of a mountain of data.' See also: Card 5 – Cognitive Search vs Knowledge Mining.
A bank wants to monitor credit card transactions for unusual behavior. Which AI workload should they use?	Anomaly Detection. WHY: Identifies outliers and deviations in time-series or event-based data. Memory Cue: 'Anomaly = abnormal blip in the pattern.' See also: Card 7 – Use cases for Anomaly Detection.
A business needs to extract key phrases and sentiment from customer reviews. Which AI workload applies?	Natural Language Processing (NLP). WHY: NLP allows systems to analyze and understand human language for tasks like sentiment analysis and key phrase extraction. Memory Cue: 'NLP = makes sense of words.' See also: Card 12 – Language Service features.
Which AI workload is used for visual tasks like classifying objects in photos or detecting faces?	Computer Vision. WHY: Enables AI to process and analyze images for tagging, object detection, OCR, and more. Memory Cue: 'Vision = seeing and understanding images.' See also: Card 14 – Computer Vision vs Custom Vision.
Which Responsible AI principle ensures systems treat all users fairly, avoiding discrimination in decisions like loan approvals?	Fairness. WHY: Ensures AI models do not propagate or amplify bias based on gender, ethnicity, or other factors. Memory Cue: 'Fair play for all.' See also: Card 20 – Fairness vs Inclusiveness.
A healthcare provider wants to ensure AI systems protect patient confidentiality. Which principle applies?	Privacy and Security. WHY: Protects personal data and ensures ethical data use. Memory Cue: 'Privacy = protect the person.' See also: Card 22 – Privacy in healthcare AI.
A company adds in-depth documentation so users understand how their AI system works and its limitations. Which principle is this?	Transparency. WHY: Makes AI systems understandable and their purpose and constraints clear. Memory Cue: 'Transparency = nothing hidden.' See also: Card 25 – Transparency vs Accountability.
An AI startup sets up strict oversight committees to monitor AI deployments for ethical and legal compliance. Which principle?	Accountability. WHY: Ensures governance structures are in place to make AI traceable and answerable. Memory Cue: 'Accountable AI = someone's responsible.' See also: Card 28 – Governance frameworks.
An AI system for public services is tested to ensure it benefits all demographics equally. Which principle applies?	Inclusiveness. WHY: Ensures AI benefits reach all users, regardless of background. Memory Cue: 'Inclusiveness = all welcome.' See also: Card 20 – Fairness vs Inclusiveness.
A company uses rigorous testing to ensure its autonomous vehicles don't malfunction in unpredictable conditions. Which principle?	Reliability and Safety. WHY: AI must function consistently under expected and unexpected scenarios. Memory Cue: 'Reliable = safe to trust.' See also: Card 30 – Safety testing in AI systems.
How is Fairness different from Inclusiveness in Responsible AI?	Fairness focuses on removing bias from AI outputs (e.g., loan approvals), while Inclusiveness ensures AI solutions benefit all societal groups. Memory Cue: 'Fair = no bias, Inclusive = everyone gains.' See also: Card 5 – Responsible AI overview.
How is Transparency different from Accountability?	Transparency ensures users understand how an AI works. Accountability ensures there is governance and oversight for AI's operation. Memory Cue: 'Transparency = clarity, Accountability = responsibility.' See also: Card 28 – Governance frameworks.